

Vaccination Rate Mini-Project: class17

Chloe J. Welch

11/26/2021

Getting started

We will start by downloading the most recently dated “Statewide COVID-19 Vaccines Administered by ZIP Code” CSV file from: <https://data.ca.gov/dataset/covid-19-vaccine-progress-dashboard-data-by-zip-code> (<https://data.ca.gov/dataset/covid-19-vaccine-progress-dashboard-data-by-zip-code>)

Be sure to move your downloaded CSV file to your project directory and then read/import into an R object called `vax`. We will use this data to answer all the questions below.

```
# Import vaccination data
vax <- read.csv("vaxdata.csv")
head(vax)
```

```
## as_of_date zip_code_tabulation_area local_health_jurisdiction county
## 1 2021-01-05 92395 San Bernardino San Bernardino
## 2 2021-01-05 93206 Kern Kern
## 3 2021-01-05 91006 Los Angeles Los Angeles
## 4 2021-01-05 91901 San Diego San Diego
## 5 2021-01-05 92230 Riverside Riverside
## 6 2021-01-05 92662 Orange Orange
## vaccine_equity_metric_quartile vem_source
## 1 1 Healthy Places Index Score
## 2 1 Healthy Places Index Score
## 3 3 Healthy Places Index Score
## 4 3 Healthy Places Index Score
## 5 1 Healthy Places Index Score
## 6 4 Healthy Places Index Score
## age12_plus_population age5_plus_population persons_fully_vaccinated
## 1 35915.3 40888 NA
## 2 1237.5 1521 NA
## 3 28742.7 31347 19
## 4 15549.8 16905 12
## 5 2320.2 2526 NA
## 6 2349.5 2397 NA
## persons_partially_vaccinated percent_of_population_fully_vaccinated
## 1 NA NA
## 2 NA NA
## 3 873 0.000606
## 4 271 0.000710
## 5 NA NA
## 6 NA NA
## percent_of_population_partially_vaccinated
## 1 NA
## 2 NA
## 3 0.027850
## 4 0.016031
## 5 NA
## 6 NA
## percent_of_population_with_1_plus_dose
## 1 NA
## 2 NA
## 3 0.028456
## 4 0.016741
## 5 NA
## 6 NA
## redacted
## 1 Information redacted in accordance with CA state privacy requirements
## 2 Information redacted in accordance with CA state privacy requirements
## 3 No
## 4 No
## 5 Information redacted in accordance with CA state privacy requirements
## 6 Information redacted in accordance with CA state privacy requirements
```

```
test <- sort(vax$as_of_date, decreasing=TRUE)
head(test)
```

```
## [1] "2021-11-23" "2021-11-23" "2021-11-23" "2021-11-23" "2021-11-23"
## [6] "2021-11-23"
```

Q1. What column details the total number of people fully vaccinated? → persons_fully_vaccinated

Q2. What column details the Zip code tabulation area?

→ zip_code_tabulation_area

Q3. What is the earliest date in this dataset?

→ 2021-01-05

Q4. What is the latest date in this dataset?

→ 2021-11-23

As we have done previously, let's call the `skim()` function from the `skimr` package to get a quick overview of this dataset:

```
library(skimr)
skimr::skim(vax)
```

Data summary

| | |
|------------------------|-------|
| Name | vax |
| Number of rows | 82908 |
| Number of columns | 14 |
| Column type frequency: | |
| character | 5 |
| numeric | 9 |
| Group variables | |
| | None |

Variable type: character

| skim_variable | n_missing | complete_rate | min | max | empty | n_unique | whitespace |
|---------------------------|-----------|---------------|-----|-----|-------|----------|------------|
| as_of_date | 0 | 1 | 10 | 10 | 0 | 47 | 0 |
| local_health_jurisdiction | 0 | 1 | 0 | 15 | 235 | 62 | 0 |
| county | 0 | 1 | 0 | 15 | 235 | 59 | 0 |
| vem_source | 0 | 1 | 15 | 26 | 0 | 3 | 0 |
| redacted | 0 | 1 | 2 | 69 | 0 | 2 | 0 |

Variable type: numeric

| skim_variable | n_missing | complete_rate | mean | sd | p0 | p25 | p50 | p75 | p100 | hist |
|--|-----------|---------------|----------|----------|-------|----------|----------|----------|----------|------|
| zip_code_tabulation_area | 0 | 1.00 | 93665.11 | 1817.39 | 90001 | 92257.75 | 93658.50 | 95380.50 | 97635.0 | |
| vaccine_equity_metric_quartile | 4089 | 0.95 | 2.44 | 1.11 | 1 | 1.00 | 2.00 | 3.00 | 4.0 | |
| age12_plus_population | 0 | 1.00 | 18895.04 | 18993.94 | 0 | 1346.95 | 13685.10 | 31756.12 | 88556.7 | |
| age5_plus_population | 0 | 1.00 | 20875.24 | 21106.04 | 0 | 1460.50 | 15364.00 | 34877.00 | 101902.0 | |
| persons_fully_vaccinated | 8355 | 0.90 | 9585.35 | 11609.12 | 11 | 516.00 | 4210.00 | 16095.00 | 71219.0 | |
| persons_partially_vaccinated | 8355 | 0.90 | 1894.87 | 2105.55 | 11 | 198.00 | 1269.00 | 2880.00 | 20159.0 | |
| percent_of_population_fully_vaccinated | 8355 | 0.90 | 0.43 | 0.27 | 0 | 0.20 | 0.44 | 0.63 | 1.0 | |
| percent_of_population_partially_vaccinated | 8355 | 0.90 | 0.10 | 0.10 | 0 | 0.06 | 0.07 | 0.11 | 1.0 | |
| percent_of_population_with_1_plus_dose | 8355 | 0.90 | 0.51 | 0.26 | 0 | 0.31 | 0.53 | 0.71 | 1.0 | |

Q5. How many numeric columns are in this dataset?

→ 9 numeric columns

Q6. Note that there are “missing values” in the dataset. How many NA values there in the `persons_fully_vaccinated` column?

```
sum(is.na(vax$persons_fully_vaccinated))
```

```
## [1] 8355
```

Q7. What percent of persons_fully_vaccinated values are missing (to 2 significant figures)?

```
sum(is.na(vax$persons_fully_vaccinated))/nrow(vax) * 100
```

```
## [1] 10.07744
```

→ To 2 significant figures = 10.

Q8. [Optional]: Why might this data be missing?

→ These values may be missing if there are records of vaccines being given (doses administered), but no specifics on the data for the people that received them.

Working with dates

One of the “character” columns of the data is as_of_date, which contains dates in the Year-Month-Day format.

Dates and times can be annoying to work with at the best of times. However, in R we have the excellent lubridate package, which can make life a lot easier. Here is a quick example to get you started:

```
library(lubridate)
```

```
##  
## Attaching package: 'lubridate'
```

```
## The following objects are masked from 'package:base':  
##  
## date, intersect, setdiff, union
```

```
today()
```

```
## [1] "2021-11-26"
```

However, if we convert our date data into a lubridate format things like this will be much easier as well as plotting time series data later on.

```
# Specify that we are using the year-month-day format  
vax$as_of_date <- ymd(vax$as_of_date)
```

Now we can do math with dates. For example: How many days have passed since the first vaccination reported in this dataset?

```
today() - vax$as_of_date[1]
```

```
## Time difference of 325 days
```

Using the last and the first date value we can now determine how many days the dataset span?

```
vax$as_of_date[nrow(vax)] - vax$as_of_date[1]
```

```
## Time difference of 322 days
```

Q9. How many days have passed since the last update of the dataset?

```
today() - vax$as_of_date[nrow(vax)]
```

```
## Time difference of 3 days
```

Q10. How many unique dates are in the dataset (i.e. how many different dates are detailed)?

```
length(unique(vax$sas_of_date))
```

```
## [1] 47
```

Working with ZIP codes

One of the numeric columns in the dataset (namely `vax$zip_code_tabulation_area`) are actually ZIP codes - a postal code used by the United States Postal Service (USPS). In R we can use the `zipcodeR` package to make working with these codes easier. For example, let's install and then load up this package and to find the centroid of the La Jolla 92037 (i.e. UC San Diego) ZIP code area.

```
library(zipcodeR)
geocode_zip('92037')
```

```
## # A tibble: 1 × 3
##   zipcode lat lng
##   <chr>   <dbl> <dbl>
## 1 92037   32.8 -117.
```

```
zip_distance('92037','92109')
```

```
##   zipcode_a zipcode_b distance
## 1      92037      92109      2.33
```

```
reverse_zipcode(c('92037', "92109") )
```

```
## # A tibble: 2 × 24
##   zipcode zipcode_type major_city post_office_city common_city_list county state
##   <chr>   <chr>         <chr>      <chr>                <blob> <chr> <chr>
## 1 92037   Standard      La Jolla    La Jolla, CA          <raw 20 B> San D... CA
## 2 92109   Standard      San Diego   San Diego, CA          <raw 21 B> San D... CA
## # ... with 17 more variables: lat <dbl>, lng <dbl>, timezone <chr>,
## #   radius_in_miles <dbl>, area_code_list <blob>, population <int>,
## #   population_density <dbl>, land_area_in_sqmi <dbl>,
## #   water_area_in_sqmi <dbl>, housing_units <int>,
## #   occupied_housing_units <int>, median_home_value <int>,
## #   median_household_income <int>, bounds_west <dbl>, bounds_east <dbl>,
## #   bounds_north <dbl>, bounds_south <dbl>
```

Optional: We can use this `reverse_zipcode()` to pull census data later on for any or all ZIP code areas we might be interested in.

```
# Pull data for all ZIP codes in the dataset
zipdata <- reverse_zipcode( vax$zip_code_tabulation_area )
head(zipdata)
```

```
## # A tibble: 6 × 24
##   zipcode zipcode_type major_city post_office_city common_city_list county state
##   <chr>    <chr>        <chr>      <chr>                <blob> <chr>  <chr>
## 1 90001    Standard    Los Angel... Los Angeles, CA      <raw 44 B> Los A... CA
## 2 90002    Standard    Los Angel... Los Angeles, CA      <raw 47 B> Los A... CA
## 3 90003    Standard    Los Angel... Los Angeles, CA      <raw 23 B> Los A... CA
## 4 90004    Standard    Los Angel... Los Angeles, CA      <raw 34 B> Los A... CA
## 5 90005    Standard    Los Angel... Los Angeles, CA      <raw 34 B> Los A... CA
## 6 90006    Standard    Los Angel... Los Angeles, CA      <raw 23 B> Los A... CA
## # ... with 17 more variables: lat <dbl>, lng <dbl>, timezone <chr>,
## #   radius_in_miles <dbl>, area_code_list <blob>, population <int>,
## #   population_density <dbl>, land_area_in_sqmi <dbl>,
## #   water_area_in_sqmi <dbl>, housing_units <int>,
## #   occupied_housing_units <int>, median_home_value <int>,
## #   median_household_income <int>, bounds_west <dbl>, bounds_east <dbl>,
## #   bounds_north <dbl>, bounds_south <dbl>
```

We could also access socioeconomic data for different ZIP code areas in a similar way if we wanted to investigate factors that might be correlated with different vaccine uptake rates.

Another informative data exploration might be to plot the various values along with the ZIP codes latitude and longitude values on a map using a package like `leaflet` or using `ggplot2` itself similar to this post. For now we will leave this as an optional extension exercise.

Focus on the San Diego area

Let's now focus in on the San Diego County area by restricting ourselves first to `vax$county == "San Diego"` entries. We have two main choices on how to do this. The first using base R the second using the `dplyr` package:

```
# Subset to San Diego county only areas
sd1 <- vax[ vax$county == "San Diego" , ]
head(sd1)
```

```
## as_of_date zip_code_tabulation_area local_health_jurisdiction county
## 4 2021-01-05 91901 San Diego San Diego
## 14 2021-01-05 91902 San Diego San Diego
## 21 2021-01-05 92011 San Diego San Diego
## 22 2021-01-05 92055 San Diego San Diego
## 25 2021-01-05 92067 San Diego San Diego
## 33 2021-01-05 92081 San Diego San Diego
## vaccine_equity_metric_quartile vem_source
## 4 3 Healthy Places Index Score
## 14 4 Healthy Places Index Score
## 21 4 Healthy Places Index Score
## 22 3 CDPH-Derived ZCTA Score
## 25 4 Healthy Places Index Score
## 33 2 Healthy Places Index Score
## age12_plus_population age5_plus_population persons_fully_vaccinated
## 4 15549.8 16905 12
## 14 16620.7 18026 22
## 21 20503.6 23247 NA
## 22 11548.0 11654 NA
## 25 6973.9 7480 11
## 33 25558.0 27632 14
## persons_partially_vaccinated percent_of_population_fully_vaccinated
## 4 271 0.000710
## 14 374 0.001220
## 21 NA NA
## 22 NA NA
## 25 241 0.001471
## 33 346 0.000507
## percent_of_population_partially_vaccinated
## 4 0.016031
## 14 0.020748
## 21 NA
## 22 NA
## 25 0.032219
## 33 0.012522
## percent_of_population_with_1_plus_dose
## 4 0.016741
## 14 0.021968
## 21 NA
## 22 NA
## 25 0.033690
## 33 0.013029
## redacted
## 4 No
## 14 No
## 21 Information redacted in accordance with CA state privacy requirements
## 22 Information redacted in accordance with CA state privacy requirements
## 25 No
## 33 No
```

```
nrow(sd1)
```

```
## [1] 5029
```

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
## filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
## intersect, setdiff, setequal, union
```

```
sd <- filter(vax, county == "San Diego")  
nrow(sd)
```

```
## [1] 5029
```

Using dplyr is often more convenient when we are subsetting across multiple criteria - for example all San Diego county areas with a population of over 10,000.

```
sd.10 <- filter(vax, county == "San Diego" &  
                age5_plus_population > 10000)  
head(sd.10)
```



```
## as_of_date zip_code_tabulation_area local_health_jurisdiction county
## 1 2021-01-05 91901 San Diego San Diego
## 2 2021-01-05 91902 San Diego San Diego
## 3 2021-01-05 92011 San Diego San Diego
## 4 2021-01-05 92055 San Diego San Diego
## 5 2021-01-05 92081 San Diego San Diego
## 6 2021-01-05 92124 San Diego San Diego
## vaccine_equity_metric_quartile vem_source
## 1 3 Healthy Places Index Score
## 2 4 Healthy Places Index Score
## 3 4 Healthy Places Index Score
## 4 3 CDPH-Derived ZCTA Score
## 5 2 Healthy Places Index Score
## 6 3 Healthy Places Index Score
## age12_plus_population age5_plus_population persons_fully_vaccinated
## 1 15549.8 16905 12
## 2 16620.7 18026 22
## 3 20503.6 23247 NA
## 4 11548.0 11654 NA
## 5 25558.0 27632 14
## 6 25422.4 29040 29
## persons_partially_vaccinated percent_of_population_fully_vaccinated
## 1 271 0.000710
## 2 374 0.001220
## 3 NA NA
## 4 NA NA
## 5 346 0.000507
## 6 575 0.000999
## percent_of_population_partially_vaccinated
## 1 0.016031
## 2 0.020748
## 3 NA
## 4 NA
## 5 0.012522
## 6 0.019800
## percent_of_population_with_1_plus_dose
## 1 0.016741
## 2 0.021968
## 3 NA
## 4 NA
## 5 0.013029
## 6 0.020799
## redacted
## 1 No
## 2 No
## 3 Information redacted in accordance with CA state privacy requirements
## 4 Information redacted in accordance with CA state privacy requirements
## 5 No
## 6 No
```

Q11. How many distinct zip codes are listed for San Diego County?

```
length(unique(sd$zip_code_tabulation_area))
```

```
## [1] 107
```

Q12. What San Diego County Zip code area has the largest 12 + Population in this dataset?

```
which.max(sd$age12_plus_population)
```

```
## [1] 60
```

```
sd$zip_code_tabulation_area[23]
```

```
## [1] 92057
```

Using `dplyr` select all San Diego “county” entries on “as_of_date” “2021-11-16” and use this for the following questions.

Q13. What is the overall average “Percent of Population Fully Vaccinated” value for all San Diego “County” as of “2021-11-09”?

```
sd_date <- filter(vax, county == "San Diego" & as_of_date == "2021-11-09")
head(sd_date)
```

```
## as_of_date zip_code_tabulation_area local_health_jurisdiction county
## 1 2021-11-09 92081 San Diego San Diego
## 2 2021-11-09 92058 San Diego San Diego
## 3 2021-11-09 91902 San Diego San Diego
## 4 2021-11-09 92140 San Diego San Diego
## 5 2021-11-09 92124 San Diego San Diego
## 6 2021-11-09 92135 San Diego San Diego
## vaccine_equity_metric_quartile vem_source
## 1 2 Healthy Places Index Score
## 2 1 Healthy Places Index Score
## 3 4 Healthy Places Index Score
## 4 NA No VEM Assigned
## 5 3 Healthy Places Index Score
## 6 NA No VEM Assigned
## age12_plus_population age5_plus_population persons_fully_vaccinated
## 1 25558.0 27632 17333
## 2 34956.0 39695 13892
## 3 16620.7 18026 13101
## 4 3747.7 3737 38
## 5 25422.4 29040 16121
## 6 635.0 635 NA
## persons_partially_vaccinated percent_of_population_fully_vaccinated
## 1 2219 0.627280
## 2 2410 0.349969
## 3 1997 0.726784
## 4 14 0.010169
## 5 2060 0.555131
## 6 NA NA
## percent_of_population_partially_vaccinated
## 1 0.080305
## 2 0.060713
## 3 0.110784
## 4 0.003746
## 5 0.070937
## 6 NA
## percent_of_population_with_1_plus_dose
## 1 0.707585
## 2 0.410682
## 3 0.837568
## 4 0.013915
## 5 0.626068
## 6 NA
## redacted
## 1 No
## 2 No
## 3 No
## 4 No
## 5 No
## 6 Information redacted in accordance with CA state privacy requirements
```

```
summary(sd_date$percent_of_population_fully_vaccinated)
```

```
## Min. 1st Qu. Median Mean 3rd Qu. Max. NA's
## 0.01017 0.60805 0.67711 0.67347 0.76257 1.00000 4
```


```
x <- sum(is.na(sd_date$percent_of_population_fully_vaccinated))
sum(na.omit(sd_date$percent_of_population_fully_vaccinated)) / (nrow(sd_date) - x) * 100
```

```
## [1] 67.34714
```

Q14. Using either ggplot or base R graphics make a summary figure that shows the distribution of Percent of Population Fully Vaccinated values as of "2021-11-09"?

```
hist(sd_date$percent_of_population_fully_vaccinated)
```

Histogram of sd_date\$percent_of_population_fully_vaccinated



Focus on UCSD/La Jolla


UC San Diego resides in the 92037 ZIP code area and is listed with an age 5+ population size of 36,144.

```
ucsd <- filter(sd, zip_code_tabulation_area=="92037")
ucsd[1,]$age5_plus_population
```

```
## [1] 36144
```

Q15. Using ggplot make a graph of the vaccination rate time course for the 92037 ZIP code area:

```
library(ggplot2)
ggplot(ucsd) +
  aes(as_of_date,
      percent_of_population_fully_vaccinated) +
  geom_point() +
  geom_line(group=1) +
  ylim(c(0,1)) +
  labs(x="Date", y="Percent Vaccinated")
```



This plot shows an initial slow roll out

in January into February (likely due to limited vaccine availability). This is followed with rapid ramp up until a clear slowing trend from June time onward. Interpretation beyond this requires context from other zip code areas to answer questions such as: is this trend representative of other areas? Are more people fully vaccinated in this area compared to others? Etc.

Comparing to similar sized areas

Let's return to the full dataset and look across every zip code area with a population at least as large as that of 92037 on as_of_date "2021-11-16".

```
# Subset to all CA areas with a population as large as 92037
vax.36 <- filter(vax, age5_plus_population > 36144 &
                  as_of_date == "2021-11-16")
head(vax.36)
```

```
## as_of_date zip_code_tabulation_area local_health_jurisdiction county
## 1 2021-11-16 92020 San Diego San Diego
## 2 2021-11-16 92563 Riverside Riverside
## 3 2021-11-16 92806 Orange Orange
## 4 2021-11-16 93291 Tulare Tulare
## 5 2021-11-16 92335 San Bernardino San Bernardino
## 6 2021-11-16 92618 Orange Orange
## vaccine_equity_metric_quartile vem_source
## 1 2 Healthy Places Index Score
## 2 3 Healthy Places Index Score
## 3 2 Healthy Places Index Score
## 4 1 Healthy Places Index Score
## 5 1 Healthy Places Index Score
## 6 4 Healthy Places Index Score
## age12_plus_population age5_plus_population persons_fully_vaccinated
## 1 49284.5 54991 35128
## 2 55897.8 63794 36051
## 3 33050.9 36739 24810
## 4 46879.7 54254 27936
## 5 79670.3 91867 49820
## 6 40348.0 44304 39695
## persons_partially_vaccinated percent_of_population_fully_vaccinated
## 1 5161 0.638795
## 2 4224 0.565116
## 3 2355 0.675304
## 4 4012 0.514911
## 5 5970 0.542306
## 6 3936 0.895969
## percent_of_population_partially_vaccinated
## 1 0.093852
## 2 0.066213
## 3 0.064101
## 4 0.073948
## 5 0.064985
## 6 0.088841
## percent_of_population_with_1_plus_dose redacted
## 1 0.732647 No
## 2 0.631329 No
## 3 0.739405 No
## 4 0.588859 No
## 5 0.607291 No
## 6 0.984810 No
```

Q16. Calculate the mean “Percent of Population Fully Vaccinated” for ZIP code areas with a population as large as 92037 (La Jolla) as_of_date “2021-11-16”. Add this as a straight horizontal line to your plot from above with the `geom_hline()` function?

```
mean(vax.36$percent_of_population_fully_vaccinated)
```


```
## [1] 0.6640413
```

Let's make a new plot:

```
ucsd$as_of_date
```

```
## [1] "2021-01-05" "2021-01-12" "2021-01-19" "2021-01-26" "2021-02-02"
## [6] "2021-02-09" "2021-02-16" "2021-02-23" "2021-03-02" "2021-03-09"
## [11] "2021-03-16" "2021-03-23" "2021-03-30" "2021-04-06" "2021-04-13"
## [16] "2021-04-20" "2021-04-27" "2021-05-04" "2021-05-11" "2021-05-18"
## [21] "2021-05-25" "2021-06-01" "2021-06-08" "2021-06-15" "2021-06-22"
## [26] "2021-06-29" "2021-07-06" "2021-07-13" "2021-07-20" "2021-07-27"
## [31] "2021-08-03" "2021-08-10" "2021-08-17" "2021-08-24" "2021-08-31"
## [36] "2021-09-07" "2021-09-14" "2021-09-21" "2021-09-28" "2021-10-05"
## [41] "2021-10-12" "2021-10-19" "2021-10-26" "2021-11-02" "2021-11-09"
## [46] "2021-11-16" "2021-11-23"
```

```
ggplot(ucsd) +
  aes(as_of_date,
    percent_of_population_fully_vaccinated) +
  geom_point() +
  geom_line(group=1) +
  ylim(c(0,1)) +
  labs(x="Date", y="Percent Vaccinated") +
  geom_hline(yintercept = 0.6640413)
```



Q17. What is the 6 number summary (Min, 1st Qu., Median, Mean, 3rd Qu., and Max) of the “Percent of Population Fully Vaccinated” values for ZIP code areas with a population as large as 92037 (La Jolla) as_of_date “2021-11-16”?

```
dat <- summary(vax.36$percent_of_population_fully_vaccinated)
head(dat)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.    Max.
## 0.3528910 0.5905170 0.6661630 0.6640413 0.7297545 1.0000000
```

Q18. Using `ggplot`, generate a histogram of this data.


```
vax.36$percent_of_population_fully_vaccinated
```

```
## [1] 0.638795 0.565116 0.675304 0.514911 0.542306 0.895969 0.610183 0.754831
## [9] 0.671440 0.685062 0.571825 0.685009 0.789706 0.758088 0.758018 0.668217
## [17] 0.533796 0.651510 0.687135 0.672017 0.705095 0.765967 0.544121 0.829537
## [25] 0.637903 0.700493 0.613925 0.628329 0.642227 0.654824 0.631281 0.682419
## [33] 0.819272 0.730465 0.573062 0.875929 0.659628 0.511820 0.706522 0.671575
## [41] 0.709131 0.526173 0.721366 0.821927 0.556778 0.536308 0.512772 0.674613
## [49] 0.690915 0.589145 0.661266 0.818517 0.703561 0.556520 0.756571 0.867086
## [57] 0.556596 0.698836 0.636051 0.777597 0.625551 0.631382 0.716066 0.588702
## [65] 0.539396 0.707375 0.718392 0.472181 0.528009 0.796706 0.644050 0.783687
## [73] 0.631097 0.956078 0.753598 0.547531 0.597529 0.480697 0.618461 0.688234
## [81] 0.610172 0.881642 0.836037 0.649573 0.565566 0.532163 0.795321 0.687544
## [89] 0.671494 0.542527 0.548829 0.573898 0.687350 0.878852 0.666937 0.841436
## [97] 0.637505 0.562748 0.677776 0.700007 0.572831 0.606870 0.553326 0.714489
## [105] 0.537228 0.750175 0.563423 0.745997 0.643037 0.749612 0.623749 0.680768
## [113] 0.767511 0.521701 0.522434 0.682254 0.523732 0.583474 0.653602 0.741917
## [121] 0.764800 0.855271 0.721193 0.701577 0.500653 0.433647 0.688582 0.631672
## [129] 0.662798 0.576452 0.601809 0.542173 0.619857 0.685675 0.716349 0.637176
## [137] 0.667082 0.780244 0.541241 0.741907 0.517657 0.685097 0.670161 0.707103
## [145] 0.767342 0.733755 0.638490 0.716598 0.759017 0.601673 0.702513 0.655895
## [153] 0.640323 0.768993 0.839498 0.684763 0.652456 0.517969 0.654527 0.654024
## [161] 0.530940 0.764964 0.742775 0.805337 0.651185 0.721270 0.614656 0.695125
## [169] 0.859553 0.728817 0.628313 0.670734 0.656297 0.764209 0.756293 0.948087
## [177] 0.690923 0.485393 0.574872 0.510786 0.610938 0.577006 0.549621 0.651296
## [185] 0.569071 0.788966 0.463319 0.623384 0.717695 0.784795 1.000000 0.658827
## [193] 0.574434 0.530863 0.654740 0.755299 0.586125 0.645119 0.436113 0.715121
## [201] 0.524522 0.657273 0.605903 0.665958 0.493331 0.771810 0.656647 0.526546
## [209] 0.603002 0.686999 0.476319 0.556440 0.668021 0.763976 0.632720 0.541915
## [217] 0.666057 0.714168 0.556318 0.743329 0.755339 0.811226 0.616480 0.813719
## [225] 0.595697 0.602320 0.653758 0.573314 0.758939 0.795904 0.620828 0.672871
## [233] 0.851894 0.584296 0.633778 0.521047 0.611774 0.784453 0.818396 0.557832
## [241] 0.549247 0.655007 0.776212 0.908691 0.842151 0.706651 0.522779 0.671247
## [249] 0.849087 0.661316 0.568585 0.552498 0.429913 0.741157 0.683580 0.518647
## [257] 0.778177 0.703912 0.530206 0.772152 0.584391 0.665908 0.820835 0.712352
## [265] 0.786689 0.704918 0.525192 0.591063 0.691787 0.775898 0.815766 0.634499
## [273] 0.686369 0.641569 0.620949 0.628331 0.749648 0.782176 0.674917 0.613536
## [281] 0.733390 0.662913 0.717260 0.708893 0.596023 0.703498 0.625933 0.750124
## [289] 0.634371 0.680762 0.669844 0.666881 0.702129 0.647831 0.844584 0.703221
## [297] 0.644819 0.602026 0.671617 0.778482 0.450352 0.696519 0.670376 0.566692
## [305] 0.666163 0.747186 0.504311 0.726094 0.798128 0.529317 0.742954 0.729044
## [313] 0.978329 0.695313 0.663354 0.603575 0.620572 0.632964 0.718851 0.611506
## [321] 0.564079 0.839156 0.821020 0.784338 0.789746 0.764371 0.546735 0.521949
## [329] 0.616061 0.667899 0.508542 0.781801 0.546277 0.628636 0.518007 0.786246
## [337] 0.676550 0.590092 0.707749 0.680295 0.590942 0.532132 0.747721 0.599918
## [345] 0.690365 0.352891 0.562106 0.468672 0.616078 0.541572 0.723034 0.591270
## [353] 0.727442 0.680905 0.653183 0.819011 0.692538 0.788603 0.454901 0.713177
## [361] 0.708313 0.489985 0.732898 0.640603 0.594066 0.619922 0.651024 0.860438
## [369] 0.671776 0.718488 0.496700 0.754363 0.589222 0.746015 0.668050 0.701483
## [377] 0.615606 0.609210 0.809691 0.544870 0.625113 0.610835 0.576914 0.694395
## [385] 0.763135 0.619944 0.649547 0.651101 0.613572 0.694361 0.839469 0.760372
## [393] 0.744376 0.680805 0.688630 0.604627 0.715254 0.524380 0.576548 0.668389
## [401] 0.818129 0.498205 0.544950 0.790066 0.524256 0.703243 0.708919 0.561631
## [409] 0.572598 0.692633 0.799640
```

```
ggplot(vax.36) +
  aes(vax.36$percent_of_population_fully_vaccinated) +
  geom_histogram()
```

```
## Warning: Use of `vax.36$percent_of_population_fully_vaccinated` is discouraged.
## Use `percent_of_population_fully_vaccinated` instead.
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



Q19. Is the 92109 and 92040 ZIP code areas above or below the average value you calculated for all these above?

```
vax %>% filter(as_of_date == "2021-11-16") %>%
  filter(zip_code_tabulation_area=="92040") %>%
  select(percent_of_population_fully_vaccinated)
```

```
## percent_of_population_fully_vaccinated
## 1 0.521047
```

```
vax %>% filter(as_of_date == "2021-11-16") %>%
  filter(zip_code_tabulation_area=="92109") %>%
  select(percent_of_population_fully_vaccinated)
```


```
## percent_of_population_fully_vaccinated
## 1 0.68863
```

→ 92109 ZIP code = above, 92040 ZIP code = below

Q20. Finally, make a time course plot of vaccination progress for all areas in the full dataset with a `age5_plus_population > 36144`.

```
vax.36.all <- filter(vax, age5_plus_population > 36144)
ggplot(vax.36.all) +
  aes(as_of_date,
      percent_of_population_fully_vaccinated,
      group=zip_code_tabulation_area) +
  geom_line(alpha=0.2, color="purple") +
  ylim(c(0,1)) +
  labs(x="Date", y="Percent Vaccinated",
       title="Vaccination Rate Across California",
       subtitle="Only areas with a population about 36K are shown") +
  geom_hline(yintercept = 0.6629812, linetype = "dashed")
```

```
## Warning: Removed 176 row(s) containing missing values (geom_path).
```

Q21. How do you feel about traveling for Thanksgiving and meeting for in-person class next week?

→ I missed our in-person class this week due to travel plans, but was able to work through the mini-project :)